

HotStuff: BFT Consensus with Linearity and Responsiveness

Yin, M., Malkhi, D., Reiter, M. K., Gueta, G. G., & Abraham, I.

PODC '19

Contents

1. Contributions
2. Prior works
3. System model
4. Review: 2-phase paradigm (PBFT, Tendermint)
5. Basic HotStuff
6. Chained HotStuff

Contributions

- **Linear view change**
 - Use a threshold signature scheme of which threshold is $2f + 1$.
- **Optimistic Responsiveness**
 - *Responsiveness* requires that a non-faulty leader can drive the protocol to consensus in time depending only on the actual message delays, independent of any known upper bound on message transmission delays
 - *Optimistic* means that *responsiveness* is required only in beneficial circumstances—here, after GST is reached.

Prior works

Protocol	<i>Correct leader</i>	Authenticator complexity			Responsiveness
		<i>Leader failure (view-change)</i>	<i>f leader failures</i>		
DLS [25]	$O(n^4)$	$O(n^4)$	$O(n^4)$		
PBFT [20]	$O(n^2)$	$O(n^3)$	$O(fn^3)$		✓
SBFT [30]	$O(n)$	$O(n^2)$	$O(fn^2)$		✓
Tendermint [15] / Casper [17]	$O(n^2)$	$O(n^2)$	$O(fn^2)$		
Tendermint* / Casper*	$O(n)$	$O(n)$	$O(fn)$		
HotStuff	$O(n)$	$O(n)$	$O(fn)$		✓

*Signatures can be combined using threshold signatures, though this optimization is not mentioned in their original works.

Table 1: Performance of selected protocols after GST.

System Model (1)

- **Security model**
 - An *adversary* coordinates Byzantine replicas and learns all internal states held by these replicas
- **Network assumption**
 - Communication is p2p, authenticated and reliable.
 - “broadcast” involves the sender.
 - Partial synchrony model.




System Model (2)

- **Threshold Signature**
 - A single public key held by all replicas, and each of the n replicas holds a distinct private key.
 - The i -th replica can use its private key to contribute a partial signature $\rho_i \leftarrow \mathbf{tsign}_i(m)$ on message m .
 - $\sigma \leftarrow \mathbf{tcombine}(m, \{\rho_i\}_{i \in I})$ on message m , where $|I| = \text{threshold} = 2f + 1$.
 - Any other replica can verify the signature using the public key and the function $\mathbf{tverify}(\sigma, m)$
- **Authenticator complexity**
 - an *authenticator* is either a partial signature or a signature.
 - Sum of the number of *authenticators* received by replica i in the protocol to reach a consensus decision.




Review: 2-phase paradigm



- **1st phase :** guarantees proposal uniqueness through the formation of a quorum certificate (QC) consisting of $(n - f)$ votes. validity condition is important.
- **2nd phase :** guarantees that the next leader can convince replicas to vote for a safe proposal, using *commit-adopt**




Review: 2-phase paradigm (cont.)



- **Commit-Adopt**: if a correct replica commits a value v , at least $2f + 1$ replicas adopt and protect v (lock the value v). At least $f + 1$ replicas are correct among those replicas, and the corresponding $f + 1$ replicas guard safety of the committed value v . Thus, they would not vote on a conflicting value unless they are shown a proof that it is safe to do so.




Review: failure cases



1. No replica has a lock, no value is committed.
2. One or a few replicas are locked on a value, but no correct replica has committed.
3. $2f + 1$ parties are locked, some correct replicas have committed but not all



Review: failure cases




- ~~1. No replica has a lock, no value is committed.~~
- ~~2. One or a few replicas are locked on a value, but no correct replica has committed.~~
- ~~3. $2f + 1$ parties are locked, some correct replicas have committed but not all~~

Review: failure cases

- In case 2, some mechanism is required to ensure that these locked correct replicas vote for a safe proposal from a leader (perhaps conflicting with a locked value) after a view-change.
 1. What does the leader learn about the status of the system (the three scenarios) at the start of a view?
 2. How does the leader convince other replicas about the status of the system (and thus to vote for its proposal)?


Review: PBFT



- Leader sends the *status certificate* containing the $2f + 1$ locks in its proposal.
- Quadratic complexity / Responsiveness




Review: Tendermint



- Avoid sending a $(2f + 1)$ -sized status certificate → linear view-change
- During Δ , the leader collects all locks and make new proposal with highest lock.
- Linear view-change / Non-responsiveness


Basic HotStuff



- *prepareQC* : qualified certificate for $2f + 1$ validity votes.
- *lockQC*: qualified certificate for $2f + 1$ locks on a proposal.




Basic HotStuff _ mitigation of case-2



- Each replica sends its own *prepareQC* (replica3 is not locked but has a *prepareQC*).
- Upon receiving $2f+1$ *prepareQCs*, leader makes new proposal with *highQC* among those *prepareQCs*.
- Locked correct replicas vote to new proposal when $\text{highQC} > \text{lockedQC}$.



Cf) HotStuff-2



- Highest qc must be found in 2-phase BFT.
- If the leader obtains a qc from the preceding view, the qc is maximal value that possibly exists in the system. In this case, it proceeds with a proposal in a responsive manner.
- Otherwise, must waits Δ . This can happen when the leader from the previous view is malicious or the previous view was before GST.

Basic HotStuff _ pseudocode(1)

Algorithm 1 Utilities (for replica r).

```

1: function  $\text{Msg}(type, node, qc)$ 
2:    $m.type \leftarrow type$ 
3:    $m.viewNumber \leftarrow curView$ 
4:    $m.node \leftarrow node$ 
5:    $m.justify \leftarrow qc$ 
6:   return  $m$ 
7: function  $\text{VOTEMSG}(type, node, qc)$ 
8:    $m \leftarrow \text{Msg}(type, node, qc)$ 
9:    $m.partialSig \leftarrow tsign_r(\langle m.type, m.viewNumber, m.node \rangle)$ 
10:  return  $m$ 
11: procedure  $\text{CREATELEAF}(parent, cmd)$ 
12:    $b.parent \leftarrow parent$ 
13:    $b.cmd \leftarrow cmd$ 
14:   return  $b$ 
15: function  $\text{QC}(V)$ 
16:    $qc.type \leftarrow m.type : m \in V$ 
17:    $qc.viewNumber \leftarrow m.viewNumber : m \in V$ 
18:    $qc.node \leftarrow m.node : m \in V$ 
19:    $qc.sig \leftarrow tcombine(\langle qc.type, qc.viewNumber, qc.node \rangle,$ 
       $\{m.partialSig \mid m \in V\})$ 
20:   return  $qc$ 
21: function  $\text{MATCHINGMSG}(m, t, v)$ 
22:   return  $(m.type = t) \wedge (m.viewNumber = v)$ 
23: function  $\text{MATCHINGQC}(qc, t, v)$ 
24:   return  $(qc.type = t) \wedge (qc.viewNumber = v)$ 
25: function  $\text{SAFENODE}(node, qc)$ 
26:   return  $(node \text{ extends from } lockedQC.node) \vee$  // safety rule
27:    $(qc.viewNumber > lockedQC.viewNumber)$  // liveness rule
  
```

Algorithm 2 Basic HotStuff protocol (for replica r).

```

1: for  $curView \leftarrow 1, 2, 3, \dots$  do
  ▷ PREPARE phase
2:   as a leader //  $r = \text{LEADER}(curView)$ 
    // we assume special NEW-VIEW messages from view 0
3:     wait for  $(n - f)$  NEW-VIEW messages:  $M \leftarrow \{m \mid \text{MATCHINGMSG}(m, \text{NEW-VIEW}, curView - 1)\}$ 
4:      $highQC \leftarrow \left( \arg \max_{m \in M} \{m.justify.viewNumber\} \right).justify$ 
5:      $curProposal \leftarrow \text{CREATELEAF}(highQC.node, \text{client's command})$ 
6:     broadcast  $\text{MSG(PREPARE, } curProposal, highQC)$ 
7:   as a replica
8:     wait for message  $m : \text{MATCHINGMSG}(m, \text{PREPARE}, curView)$  from  $\text{LEADER}(curView)$ 
9:     if  $m.node$  extends from  $m.justify.node \wedge$ 
         $\text{SAFENODE}(m.node, m.justify)$  then
10:       send  $\text{VOTEMSG}(\text{PREPARE}, m.node, \perp)$  to  $\text{LEADER}(curView)$ 
  ▷ PRE-COMMIT phase
11:   as a leader
12:     wait for  $(n - f)$  votes:  $V \leftarrow \{v \mid \text{MATCHINGMSG}(v, \text{PREPARE}, curView)\}$ 
13:      $prepareQC \leftarrow \text{QC}(V)$ 
14:     broadcast  $\text{MSG(PRE-COMMIT, } \perp, prepareQC)$ 
15:   as a replica
16:     wait for message  $m : \text{MATCHINGQC}(m.justify, \text{PREPARE}, curView)$  from  $\text{LEADER}(curView)$ 
17:      $prepareQC \leftarrow m.justify$ 
18:     send  $\text{VOTEMSG}(\text{PRE-COMMIT}, m.justify.node, \perp)$  to  $\text{LEADER}(curView)$ 
  
```

Basic HotStuff _ pseudocode(2)

Algorithm 1 Utilities (for replica r).

```

1: function Msg( $type, node, qc$ )
2:    $m.type \leftarrow type$ 
3:    $m.viewNumber \leftarrow curView$ 
4:    $m.node \leftarrow node$ 
5:    $m.justify \leftarrow qc$ 
6:   return  $m$ 
7: function VOTEMSG( $type, node, qc$ )
8:    $m \leftarrow \text{Msg}(type, node, qc)$ 
9:    $m.partialSig \leftarrow tsign_r(\langle m.type, m.viewNumber, m.node \rangle)$ 
10:  return  $m$ 
11: procedure CREATELEAF( $parent, cmd$ )
12:    $b.parent \leftarrow parent$ 
13:    $b.cmd \leftarrow cmd$ 
14:   return  $b$ 
15: function QC( $V$ )
16:    $qc.type \leftarrow m.type : m \in V$ 
17:    $qc.viewNumber \leftarrow m.viewNumber : m \in V$ 
18:    $qc.node \leftarrow m.node : m \in V$ 
19:    $qc.sig \leftarrow tcombine(\langle qc.type, qc.viewNumber, qc.node \rangle,$ 
         $\{m.partialSig \mid m \in V\})$ 
20:   return  $qc$ 
21: function MATCHINGMSG( $m, t, v$ )
22:   return  $(m.type = t) \wedge (m.viewNumber = v)$ 
23: function MATCHINGQC( $qc, t, v$ )
24:   return  $(qc.type = t) \wedge (qc.viewNumber = v)$ 
25: function SAFENODE( $node, qc$ )
26:   return  $(node \text{ extends from } lockedQC.node) \vee$  // safety rule
27:    $(qc.viewNumber > lockedQC.viewNumber)$  // liveness rule
  
```

▷ **COMMIT phase**

19: **as a leader**

20: wait for $(n - f)$ votes: $V \leftarrow \{v \mid \text{MATCHINGMSG}(v, \text{PRE-COMMIT}, curView)\}$

21: $precommitQC \leftarrow QC(V)$

22: broadcast $\text{Msg}(\text{COMMIT}, \perp, precommitQC)$

23: **as a replica**

24: wait for message $m : \text{MATCHINGQC}(m.justify, \text{PRE-COMMIT}, curView)$ from $\text{LEADER}(curView)$

25: $lockedQC \leftarrow m.justify$

26: send $\text{VOTEMSG}(\text{COMMIT}, m.justify.node, \perp)$ to $\text{LEADER}(curView)$

▷ **DECIDE phase**

27: **as a leader**

28: wait for $(n - f)$ votes: $V \leftarrow \{v \mid \text{MATCHINGMSG}(v, \text{COMMIT}, curView)\}$

29: $commitQC \leftarrow QC(V)$

30: broadcast $\text{Msg}(\text{DECIDE}, \perp, commitQC)$

31: **as a replica**

32: wait for message m from $\text{LEADER}(curView)$

33: wait for message $m : \text{MATCHINGQC}(m.justify, \text{COMMIT}, curView)$ from $\text{LEADER}(curView)$


34: execute new commands through $m.justify.node$, respond to clients

▷ **Finally**

35: NEXTVIEW interrupt: goto this line if $\text{NEXTVIEW}(curView)$ is called during “wait for” in any phase

36: send $\text{Msg}(\text{NEW-VIEW}, \perp, prepareQC)$ to $\text{LEADER}(curView + 1)$

Chained HotStuff



Chained HotStuff _ pseudocode

```

6: for  $curView \leftarrow 1, 2, 3, \dots$  do
    ▷ GENERIC phase
    7:   as a leader //  $r = \text{LEADER}(curView)$ 
        //  $M$  is the set of messages collected at the end of previous view by the leader of this view
        8:      $high\ QC \leftarrow \left( \arg \max_{m \in M} \{m.\text{justify}.viewNumber\} \right).justify$ 
        9:     if  $high\ QC.\text{viewNumber} > generic\ QC.\text{viewNumber}$  then  $generic\ QC \leftarrow high\ QC$ 
        10:     $curProposal \leftarrow \text{CREATELEAF}(generic\ QC.\text{node}, \text{client's command}, generic\ QC)$ 
            // PREPARE phase
        11:    broadcast  $\text{MSG(GENERIC, } curProposal, \perp)$ 
        12:   as a replica
        13:     wait for message  $m : \text{MATCHINGMSG}(m, \text{GENERIC}, curView)$  from  $\text{LEADER}(curView)$ 
        14:      $b^* \leftarrow m.\text{node}; b'' \leftarrow b^*.justify.\text{node}; b' \leftarrow b''.justify.\text{node}; b \leftarrow b'.justify.\text{node}$ 
        15:     if  $\text{SAFENODE}(b^*, b^*.justify)$  then
        16:       send  $\text{VOTEMSG(GENERIC, } b^*, \perp)$  to  $\text{LEADER}(curView + 1)$ 
            // start PRE-COMMIT phase on  $b^*$ 's parent
        17:       if  $b^*.parent = b''$  then
        18:          $generic\ QC \leftarrow b^*.justify$ 
            // start COMMIT phase on  $b^*$ 's grandparent
        19:         if  $(b^*.parent = b'') \wedge (b''.parent = b')$  then
        20:            $locked\ QC \leftarrow b''.justify$ 
            // start DECIDE phase on  $b^*$ 's great-grandparent
        21:           if  $(b^*.parent = b'') \wedge (b''.parent = b') \wedge (b'.parent = b)$  then
        22:             execute new commands through  $b$ , respond to clients
        23:   as the next leader
        24:     wait for all messages:  $M \leftarrow \{m \mid \text{MATCHINGMSG}(m, \text{GENERIC}, curView)\}$ 
        25:     until there are  $(n - f)$  votes:  $V \leftarrow \{v \mid v.\text{partialSig} \neq \perp \wedge v \in M\}$ 
        26:      $generic\ QC \leftarrow \text{QC}(V)$ 
        ▷ Finally
        27:      $\text{NEXTVIEW interrupt: goto this line if NEXTVIEW(curView) is called during "wait for" in any phase}$ 
        27:     send  $\text{MSG(GENERIC, } \perp, generic\ QC)$  to  $\text{LEADER}(curView + 1)$ 

```

Evaluation _ setup

- EC2, c5.4xlarge (16 vCPUs supported by Intel Xeon Platinum 8000 processors @ 3.4GHz)
- The maximum TCP bandwidth measured by iperf was around 1.2 Gigabytes per second.
- The network latency between two machines was less than 1 ms. (LAN?)

Evaluation _ base performance




Figure 4: Throughput vs. latency with different choices of batch size, 4 replicas, 0/0 payload.






Figure 5: Throughput vs. latency with different choices of payload size, 4 replicas, batch size of 400.

Evaluation _ scalability



(a) Throughput



(b) Latency

Figure 6: Scalability with 0/0 payload, batch size of 400.

Evaluation _ scalability




Figure 7: Scalability for 128/128 payload or 1024/1024 payload, with batch size of 400.




Figure 8: Scalability for inter-replica latency 5ms ± 0.5ms or 10ms ± 1.0ms, with 0/0 payload, batch size of 400.

Evaluation _ view change




Figure 10: Number of extra authenticators used for each BFT-SMaRt view change.